

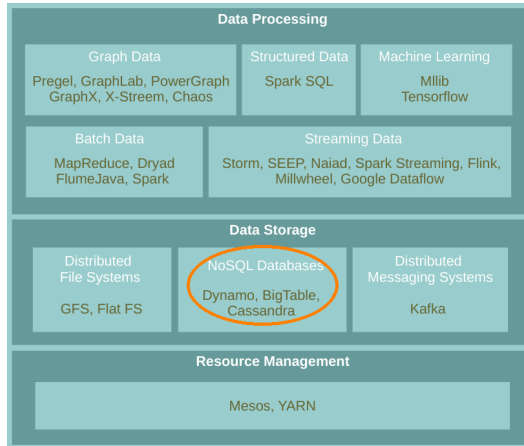


NoSQL Databases

Amir H. Payberah
payberah@kth.se
2025-09-01




Where Are We?



Database and Database Management System

- ▶ **Database:** an **organized** collection of **data**.
- ▶ **Database Management System (DBMS):** a **software** to **capture** and **analyze** data.



SQL vs. NoSQL Databases

Relational SQL Databases

- ▶ The **dominant** technology for storing **structured** data in web and business applications.
- ▶ **SQL** is good
 - **Rich** language and toolset
 - **Easy** to use and integrate
 - Many **vendors**
- ▶ They promise: **ACID**





ACID Properties

► Atomicity

- All included statements in a transaction are either **executed** or the **whole** transaction is **aborted** without affecting the database.



ACID Properties

▶ Atomicity

- All included statements in a transaction are either **executed** or the **whole** transaction is **aborted** without affecting the database.

▶ Consistency

- A database is in a **consistent** state before and after a transaction.



ACID Properties

▶ Atomicity

- All included statements in a transaction are either **executed** or the **whole** transaction is **aborted** without affecting the database.

▶ Consistency

- A database is in a **consistent** state before and after a transaction.

▶ Isolation

- Transactions can not see **uncommitted changes** in the database.



ACID Properties

▶ Atomicity

- All included statements in a transaction are either **executed** or the **whole** transaction is **aborted** without affecting the database.

▶ Consistency

- A database is in a **consistent** state before and after a transaction.

▶ Isolation

- Transactions can not see **uncommitted changes** in the database.

▶ Durability

- Changes are written to a **disk** before a database commits a transaction so that committed data cannot be lost through a power **failure**.

SQL Databases Challenges

- ▶ Web-based applications caused spikes.
 - Internet-scale data size
 - High read-write rates
 - Frequent schema changes



SQL Databases Challenges

- ▶ Web-based applications caused spikes.
 - Internet-scale data size
 - High read-write rates
 - Frequent schema changes
- ▶ RDBMS were **not** designed to be **distributed**.



► Avoids:

- Overhead of ACID properties
- Complexity of SQL query

► Provides:

- Scalability
- Easy and frequent changes to DB
- Large data volumes



Availability vs. Consistency



Availability

- ▶ Replicating data to improve the availability of data.

Availability


- ▶ Replicating data to improve the availability of data.
- ▶ Data replication
 - Storing data in more than one site or node



Consistency

► Strong consistency

- After an update completes, any subsequent access will return the **updated value**.



Consistency


► Strong consistency

- After an update completes, any subsequent access will return the **updated value**.



► Eventual consistency

- Does **not guarantee** that subsequent accesses will return the **updated value**.
- Inconsistency window**.
- If no new updates are made to the object, **eventually** all accesses will return the last updated value.





Availability vs. Consistency

- ▶ The large-scale applications have to be **reliable**: consistency, availability, partition tolerance



Availability vs. Consistency

- ▶ The large-scale applications have to be **reliable**: **consistency**, **availability**, **partition tolerance**
- ▶ Achieving **ACID** properties on large-scale applications is **challenging**.



Availability vs. Consistency

- ▶ The large-scale applications have to be **reliable**: **consistency**, **availability**, **partition tolerance**
- ▶ Achieving **ACID** properties on large-scale applications is **challenging**.
- ▶ **CAP** theorem



Go to www.menti.com, and use the code 5587 7038


► A Must for Distributed Databases?

1. Consistency
2. Availability
3. Partition Tolerance

CAP Theorem

► Consistency

- Consistent state of data after the execution of an operation.




CAP Theorem

► Consistency

- Consistent state of data after the execution of an operation.

► Availability

- Clients can always read and write data.



CAP Theorem

► Consistency


- Consistent state of data after the execution of an operation.

► Availability

- Clients can always read and write data.

► Partition Tolerance

- Continue the operation in the presence of network partitions.



CAP Theorem

► Consistency

- Consistent state of data after the execution of an operation.


► Availability

- Clients can always read and write data.

► Partition Tolerance


- Continue the operation in the presence of network partitions.

► You can choose only two!



NoSQL Data Models

NoSQL Data Models



[<http://highlyscalable.wordpress.com/2012/03/01/nosql-data-modeling-techniques>]



Key-Value Data Model

- ▶ Collection of **key/value** pairs.
- ▶ **Ordered** Key-Value: processing over **key ranges**.
- ▶ **Dynamo, Scalaris, Voldemort, Riak, ...**

Column-Oriented Data Model

- ▶ Similar to a **key/value** store, but the **value** can have multiple **attributes** (Columns).
- ▶ **Column**: a set of data **values** of a particular **type**.
- ▶ Store and process data by **column** instead of **row**.
- ▶ **BigTable**, **Hbase**, **Cassandra**, ...



Document Data Model

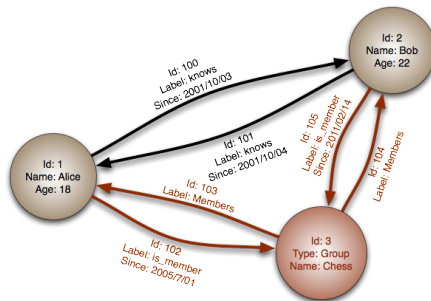
- ▶ Similar to a **column-oriented** store, but values can have **complex documents**.
- ▶ Flexible schema (XML, YAML, JSON, and BSON).
- ▶ **CouchDB, MongoDB, ...**

```
{
  FirstName: "Bob",
  Address: "5 Oak St.",
  Hobby: "sailing"
}

{
  FirstName: "Jonathan",
  Address: "15 Wanamassa Point Road",
  Children: [
    {Name: "Michael", Age: 10},
    {Name: "Jennifer", Age: 8},
  ]
}
```

Graph Data Model

- Uses **graph** structures with **nodes**, **edges**, and **properties** to represent and store data.
- Neo4J, InfoGrid, ...



[http://en.wikipedia.org/wiki/Graph_database]

BigTable

BigTable

- ▶ Lots of (semi-)structured data at Google.
 - URLs, per-user data, geographical locations, ...
- ▶ Distributed multi-level map
- ▶ CAP: strong consistency and partition tolerance



Data Model



Data Model (1/5)

- ▶ Table
- ▶ Distributed multi-dimensional sparse [map](#)



Data Model (2/5)


- ▶ Rows
- ▶ Every read or write in a row is atomic.
- ▶ Rows sorted in lexicographical order.



- ▶ Column
- ▶ The basic unit of data access.
- ▶ Column families: group of (the same type) column keys.
- ▶ Column key naming: family:qualifier

Data Model (4/5)

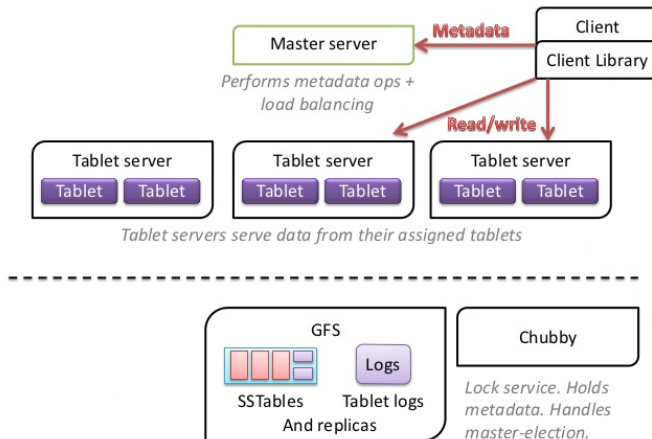
- ▶ Timestamp
- ▶ Each column value may contain multiple **versions**.



-
- The diagram illustrates two scenarios of URL normalization. In the top scenario, a URL with a missing TLD is normalized by adding a default TLD. In the bottom scenario, multiple URLs with missing TLDs are normalized to a single default TLD.
- Top Scenario:**
- Input: "content:" "anchor:cnnsi.com" "anchor:my.look.ca"
 - Input: "com.aaa"
 - Input: "com.cnn.www"
 - Output: "com.cnn.www.tech"
- Bottom Scenario:**
- Input: "content:" "anchor:cnnsi.com" "anchor:my.look.ca"
 - Input: "com.weather"
 - Input: "com.wikipedia"
 - Input: "com.zoom"

System Architecture


BigTable System Structure




[<https://www.slideshare.net/GrishaWeintraub/cap-28353551>]

Main Components

- ▶ Master
- ▶ Tablet server
- ▶ Client library




- Assigns tablets to tablet server.




Master

- ▶ Assigns tablets to tablet server.
- ▶ Balances tablet server load.




Master

- ▶ Assigns tablets to tablet server.
- ▶ Balances tablet server load.
- ▶ Garbage collection of unneeded files in GFS.




Master

- ▶ Assigns tablets to tablet server.
- ▶ Balances tablet server load.
- ▶ Garbage collection of unneeded files in GFS.
- ▶ Handles schema changes, e.g., table and column family creations




Tablet Server

- Can be added or removed dynamically.




Tablet Server

- ▶ Can be **added** or **removed** **dynamically**.
- ▶ Each **manages** a set of tablets (typically 10-1000 tablets/server).




Tablet Server

- ▶ Can be **added** or **removed** **dynamically**.
- ▶ Each **manages** a set of tablets (typically 10-1000 tablets/server).
- ▶ Handles **read/write** requests to tablets.




Tablet Server

- ▶ Can be **added** or **removed** **dynamically**.
- ▶ Each **manages** a set of tablets (typically 10-1000 tablets/server).
- ▶ Handles **read/write** requests to tablets.
- ▶ **Splits** tablets when too large.




Client Library

- **Library** that is linked into every client.




Client Library

- ▶ **Library** that is linked into every client.
- ▶ Client **data** **does not move** though the **master**.




Client Library

- ▶ **Library** that is linked into every client.
- ▶ Client **data** **does not move** though the **master**.
- ▶ Clients communicate **directly** with **tablet servers** for **reads/writes**.




Building Blocks

- ▶ The building blocks for the BigTable are:
 - Google File System (GFS)
 - Chubby
 - SSTable




Google File System (GFS)

- ▶ Large-scale distributed file system.
- ▶ Store log and data files.




Chubby Lock Service

- Ensure there is only **one active master**.




Chubby Lock Service

- ▶ Ensure there is only **one active master**.
- ▶ Store **bootstrap location** of BigTable data.




Chubby Lock Service

- ▶ Ensure there is only **one active master**.
- ▶ Store **bootstrap location** of BigTable data.
- ▶ **Discover** tablet servers.





Chubby Lock Service



- ▶ Ensure there is only **one active master**.
- ▶ Store **bootstrap location** of BigTable data.
- ▶ **Discover** tablet servers.
- ▶ Store BigTable **schema** information and **access control lists**.





- **SSTable** file format used internally to store BigTable data.





- ▶ **SSTable** file format used internally to store BigTable data.
- ▶ Chunks of **data** plus a **block index**.



- ▶ **SSTable** file format used internally to store BigTable data.
- ▶ Chunks of data plus a block index.
- ▶ **Immutable**, sorted file of key-value pairs.



- ▶ **SSTable** file format used internally to store BigTable data.
- ▶ Chunks of **data** plus a **block index**.
- ▶ **Immutable**, sorted file of **key-value** pairs.
- ▶ Each SSTable is stored in a **GFS file**.





Go to www.menti.com, and use the code 7595 3158

► Who takes care of data replication in BigTable?

1. The Master
2. GFS
3. The Chubby
4. Table Servers



Tablet Assignment

- ▶ 1 tablet \rightarrow 1 tablet server.



Tablet Assignment

- ▶ 1 tablet → 1 tablet server.
- ▶ Master uses **Chubby** to keep tracks of **live** tablet serves and **unassigned** tablets.
 - When a **tablet server starts**, it creates and acquires an **exclusive lock** in Chubby.



Tablet Assignment

- ▶ 1 tablet → 1 tablet server.
- ▶ Master uses **Chubby** to keep tracks of **live** tablet serves and **unassigned tablets**.
 - When a **tablet server starts**, it creates and acquires an **exclusive lock** in Chubby.
- ▶ Master detects the **status of the lock of each tablet server** by checking periodically.




Tablet Assignment

- ▶ 1 tablet → 1 tablet server.
- ▶ Master uses Chubby to keep tracks of live tablet serves and unassigned tablets.
 - When a tablet server starts, it creates and acquires an exclusive lock in Chubby.
- ▶ Master detects the status of the lock of each tablet server by checking periodically.
- ▶ Master is responsible for finding when tablet server is no longer serving its tablets and reassigning those tablets as soon as possible.


Finding a Tablet

- ▶ Three-level hierarchy.




Finding a Tablet

- ▶ Three-level hierarchy.
- ▶ The first level is a file stored in Chubby that contains the location of the root tablet.




Finding a Tablet

- ▶ Three-level hierarchy.
- ▶ The first level is a file stored in Chubby that contains the location of the root tablet.
- ▶ Root tablet contains location of all tablets in a special METADATA table.




Finding a Tablet

- ▶ Three-level hierarchy.
- ▶ The first level is a file stored in Chubby that contains the location of the root tablet.
- ▶ Root tablet contains location of all tablets in a special METADATA table.
- ▶ METADATA table contains location of each tablet under a row.



Finding a Tablet

- ▶ Three-level hierarchy.
- ▶ The first level is a file stored in Chubby that contains the location of the root tablet.
- ▶ Root tablet contains location of all tablets in a special METADATA table.
- ▶ METADATA table contains location of each tablet under a row.
- ▶ The client library caches tablet locations.





Loading Tablets

- ▶ To load a tablet, a tablet server does the following:



Loading Tablets

- ▶ To load a tablet, a tablet server does the following:
- ▶ Finds location of tablet through its METADATA.
 - Metadata for a tablet includes list of SSTables and set of redo points.



Loading Tablets

- ▶ To load a tablet, a tablet server does the following:
- ▶ Finds location of tablet through its METADATA.
 - Metadata for a tablet includes list of SSTables and set of redo points.
- ▶ Read SSTables index blocks into memory.




Loading Tablets

- ▶ To load a tablet, a tablet server does the following:
- ▶ Finds location of tablet through its METADATA.
 - Metadata for a tablet includes list of SSTables and set of redo points.
- ▶ Read SSTables index blocks into memory.
- ▶ Read the commit log since the redo point and reconstructs the memtable.


Tablet Serving (1/2)


- Updates committed to a [commit log](#).



Tablet Serving (1/2)

- Updates committed to a [commit log](#).
- Recently committed updates are stored in [memory](#) - [memtable](#)







Tablet Serving (2/2)

- ▶ Strong consistency
 - Only one tablet server is responsible for a given piece of data.
 - Replication is handled on the GFS layer.



Tablet Serving (2/2)

- ▶ Strong consistency
 - Only **one tablet server** is responsible for a given piece of data.
 - **Replication** is handled on the **GFS** layer.
- ▶ Trade-off with **availability**
 - If a tablet server fails, its portion of data is **temporarily unavailable** until a new server is assigned.

BigTable vs. HBase

BigTable	HBase
GFS	HDFS
Tablet Server	Region Server
SSTable	StoreFile
Memtable	MemStore
Chubby	ZooKeeper

Cassandra

- ▶ A column-oriented database
- ▶ It was created for Facebook and was later open sourced
- ▶ CAP: availability and partition tolerance





Borrowed From BigTable

- ▶ Data model: **column oriented**
 - **Keyspaces** (similar to the schema in a relational database), **tables**, and **columns**.




Borrowed From BigTable

- ▶ Data model: **column oriented**
 - **Keyspaces** (similar to the schema in a relational database), **tables**, and **columns**.
- ▶ **SSTable** disk storage
 - Append-only commit log
 - Memtable (buffering and sorting)
 - Immutable sstable files

Data Partitioning (1/2)

- ▶ **Key/value**, where values are stored as **objects**.
- ▶ If size of data **exceeds the capacity** of a single machine: **partitioning**



Data Partitioning (1/2)

- ▶ **Key/value**, where values are stored as **objects**.
- ▶ If size of data **exceeds the capacity** of a single machine: **partitioning**
- ▶ **Consistent hashing** for partitioning.






Data Partitioning (2/2)

- ▶ Consistent hashing.
- ▶ Hash both data and node ids using the same hash function in a same id space.
- ▶ $\text{partition} = \text{hash}(d) \bmod n$, d : data, n : the size of the id space

Data Partitioning (2/2)


- ▶ Consistent hashing.
- ▶ Hash both data and node ids using the same hash function in a same id space.
- ▶ $\text{partition} = \text{hash}(d) \bmod n$, d : data, n : the size of the id space

```
id space = [0, 15], n = 16  
hash("Fatemeh") = 12  
hash("Ahmad") = 2  
hash("Seif") = 9  
hash("Jim") = 14  
hash("Sverker") = 4
```



Replication


- To achieve high **availability** and **durability**, data should be **replicated** on multiple nodes.



Go to www.menti.com, and use the code 8650 0544

► $\text{hash}(\text{"A"}) = 15$ and $\text{hash}(\text{"B"}) = 5$?

1. 0 and 5
2. 15 and 5
3. 0 and 6







NoSQL Through a Data Feminism Lens (1/2)

- ▶ Who holds power?



NoSQL Through a Data Feminism Lens (1/2)

► Who holds power?

- Corporate-managed (Google, AWS, Microsoft).
- Users depend on vendors for infrastructure, pricing, and availability.



NoSQL Through a Data Feminism Lens (1/2)

- ▶ Who holds power?
 - Corporate-managed (Google, AWS, Microsoft).
 - Users depend on vendors for infrastructure, pricing, and availability.

- ▶ Who gets excluded?



NoSQL Through a Data Feminism Lens (1/2)

▶ Who holds power?

- Corporate-managed (Google, AWS, Microsoft).
- Users depend on vendors for infrastructure, pricing, and availability.

▶ Who gets excluded?

- High costs, proprietary APIs, and limited offline capabilities exclude small orgs, NGOs, and low-connectivity regions.



NoSQL Through a Data Feminism Lens (2/2)

- ▶ Who bears the **environmental cost**?



NoSQL Through a Data Feminism Lens (2/2)

- ▶ Who bears the **environmental cost**?
 - Always-on clusters = high energy and water use
 - Carbon costs externalized to hosting regions



NoSQL Through a Data Feminism Lens (2/2)

- ▶ Who bears the **environmental cost**?
 - Always-on clusters = high energy and water use
 - Carbon costs externalized to hosting regions

- ▶ Whose **labor** is **invisible**?



NoSQL Through a Data Feminism Lens (2/2)

- ▶ Who bears the **environmental cost**?
 - Always-on clusters = high energy and water use
 - Carbon costs externalized to hosting regions

- ▶ Whose **labor** is **invisible**?
 - Hardware built from mined minerals under unsafe conditions
 - Outsourced teams handle maintenance and data prep.



Questions

- ▶ If BigTable or Cassandra were reimagined for equity, inclusion, and sustainability, what would change in their [architecture, governance and replication strategy](#)?



Possible Answers

► Architecture



Possible Answers

► Architecture

- Allow flexible representation (e.g., non-binary genders, multilingual text).
- Context-aware sharding, i.e., data sharding designed to respect community.
- Built-in mechanisms for deletion.



Possible Answers

► Architecture

- Allow flexible representation (e.g., non-binary genders, multilingual text).
- Context-aware sharding, i.e., data sharding designed to respect community.
- Built-in mechanisms for deletion.

► Governance



Possible Answers

► Architecture

- Allow flexible representation (e.g., non-binary genders, multilingual text).
- Context-aware sharding, i.e., data sharding designed to respect community.
- Built-in mechanisms for deletion.

► Governance

- Shared authority on table design, retention, and access policies.
- Transparent query auditing, i.e., logs visible to users and communities, not just operators.
- Publish energy and carbon use alongside system metrics.

► Architecture

- Allow flexible representation (e.g., non-binary genders, multilingual text).
- Context-aware sharding, i.e., data sharding designed to respect community.
- Built-in mechanisms for deletion.

► Governance

- Shared authority on table design, retention, and access policies.
- Transparent query auditing, i.e., logs visible to users and communities, not just operators.
- Publish energy and carbon use alongside system metrics.

► Replication Strategy

► Architecture

- Allow flexible representation (e.g., non-binary genders, multilingual text).
- Context-aware sharding, i.e., data sharding designed to respect community.
- Built-in mechanisms for deletion.

► Governance

- Shared authority on table design, retention, and access policies.
- Transparent query auditing, i.e., logs visible to users and communities, not just operators.
- Publish energy and carbon use alongside system metrics.

► Replication Strategy

- Stricter consistency for sensitive health data, lighter for public archives.
- Adapt replication based on renewable energy availability.



Feminist-Aligned Alternative Systems

► CouchDB

- Decentralized and offline-first
- Local replicas support underserved regions and data sovereignty



Feminist-Aligned Alternative Systems

▶ CouchDB

- Decentralized and offline-first
- Local replicas support underserved regions and data sovereignty

▶ OrbitDB

- Peer-to-peer, IPFS-based
- No central authority, runs in browsers or low-power devices



Feminist-Aligned Alternative Systems

▶ CouchDB

- Decentralized and offline-first
- Local replicas support underserved regions and data sovereignty

▶ OrbitDB

- Peer-to-peer, IPFS-based
- No central authority, runs in browsers or low-power devices

▶ GunDB

- Decentralized graph database with user-owned, encrypted data
- Works offline-first



Questions

- ▶ What trade-offs might we face when choosing decentralized or community-controlled NoSQL over BigTable or Cassandra?



Possible Answers

- ▶ **Performance vs. Autonomy:** lower throughput/latency than hyperscale systems, but more local control.



Possible Answers

- ▶ **Performance vs. Autonomy:** lower throughput/latency than hyperscale systems, but more local control.
- ▶ **Simplicity vs. Flexibility:** community systems may lack advanced features or tooling, but are easier to adapt to local needs.

Possible Answers

- ▶ **Performance vs. Autonomy:** lower throughput/latency than hyperscale systems, but more local control.
- ▶ **Simplicity vs. Flexibility:** community systems may lack advanced features or tooling, but are easier to adapt to local needs.
- ▶ **Resilience vs. Cost:** fewer replicas — $>$ lower energy/cost, but higher risk of data loss.

Possible Answers

- ▶ **Performance vs. Autonomy:** lower throughput/latency than hyperscale systems, but more local control.
- ▶ **Simplicity vs. Flexibility:** community systems may lack advanced features or tooling, but are easier to adapt to local needs.
- ▶ **Resilience vs. Cost:** fewer replicas — $>$ lower energy/cost, but higher risk of data loss.
- ▶ **Global Reach vs. Sovereignty:** corporate clouds offer worldwide availability, while community NoSQL prioritizes jurisdictional/local storage.

Possible Answers

- ▶ **Performance vs. Autonomy:** lower throughput/latency than hyperscale systems, but more local control.
- ▶ **Simplicity vs. Flexibility:** community systems may lack advanced features or tooling, but are easier to adapt to local needs.
- ▶ **Resilience vs. Cost:** fewer replicas — $>$ lower energy/cost, but higher risk of data loss.
- ▶ **Global Reach vs. Sovereignty:** corporate clouds offer worldwide availability, while community NoSQL prioritizes jurisdictional/local storage.
- ▶ **Support vs. Self-reliance:** bigTech = professional 24/7 support; community systems rely on shared governance and volunteer labor.

Summary



Summary

- ▶ NoSQL data models: key-value, column-oriented, document-oriented, graph-based
- ▶ CAP (Consistency vs. Availability)



Summary

- ▶ NoSQL data models: key-value, column-oriented, document-oriented, graph-based
- ▶ CAP (Consistency vs. Availability)
- ▶ BigTable
 - Column-oriented
 - Basic components: GFS, SSTable, Chubby



Summary

- ▶ NoSQL data models: key-value, column-oriented, document-oriented, graph-based
- ▶ CAP (Consistency vs. Availability)
- ▶ BigTable
 - Column-oriented
 - Basic components: GFS, SSTable, Chubby
- ▶ Cassandra
 - Column-oriented
 - Consistency hashing



Summary

- ▶ NoSQL data models: key-value, column-oriented, document-oriented, graph-based
- ▶ CAP (Consistency vs. Availability)
- ▶ BigTable
 - Column-oriented
 - Basic components: GFS, SSTable, Chubby
- ▶ Cassandra
 - Column-oriented
 - Consistency hashing
- ▶ Alternative systems, e.g., CouchDB, OrbitDB, GunDB

- ▶ F. Chang et al., Bigtable: A distributed storage system for structured data, ACM Transactions on Computer Systems (TOCS) 26.2, 2008.
- ▶ A. Lakshman et al., Cassandra: a decentralized structured storage system, ACM SIGOPS Operating Systems Review 44.2, 2010.
- ▶ J.C. Anderson et al., CouchDB: The Definitive Guide, O'Reilly Media, 2010.
- ▶ OrbitDB: <https://orbitdb.org>
- ▶ GunDB: <https://gun.eco>

Questions?