

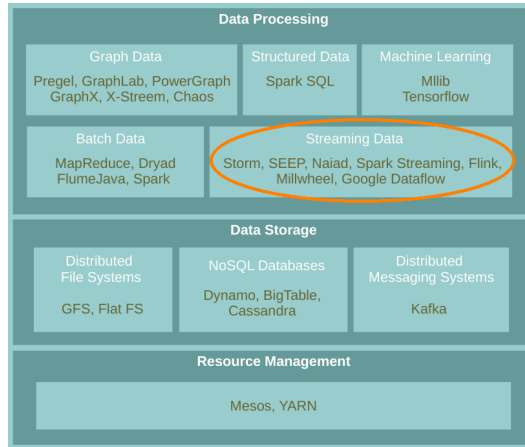


Data Stream Processing

Amir H. Payberah
payberah@kth.se
2025-09-16

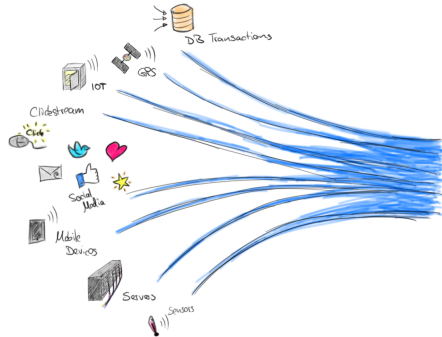


Where Are We?



Stream Processing

- ▶ **Stream processing** is the **real-time** computation of **continuously** incoming data.
- ▶ The **input data** is **unbounded**: a **series of events**, no predetermined **beginning or end**.





Streaming Data

- ▶ Data stream is unbound data, which is broken into a sequence of individual tuples.
- ▶ A data tuple is the atomic data item in a data stream.



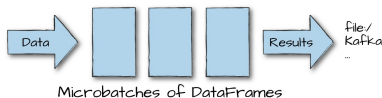
Streaming Data

- ▶ Data stream is unbound data, which is broken into a sequence of individual tuples.
- ▶ A data tuple is the atomic data item in a data stream.
- ▶ Can be structured, semi-structured, and unstructured.

Streaming Processing Patterns

► Micro-batch systems

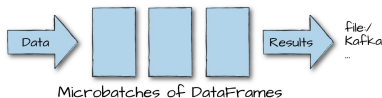
- Batch engines
- Slicing up the unbounded data into a sets of bounded data, then process each batch.



Streaming Processing Patterns

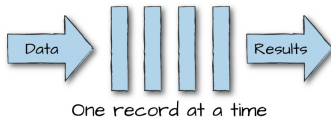
► Micro-batch systems

- Batch engines
- Slicing up the unbounded data into a sets of bounded data, then process each batch.



► Continuous processing-based systems

- Each node in the system continually listens to messages from other nodes and outputs new updates to its child nodes.



Event and Processing Time



Event Time vs. Processing Time

- ▶ **Event time**: the time at which events **actually occurred**.
 - Timestamps inserted into each record **at the source**.
- ▶ **Processing time**: the time when the record is **received at the streaming application**.



Windowing and Triggering

- ▶ **Windowing:** dividing a continuous stream of data into discrete chunks (**windows**).



Windowing and Triggering

- ▶ **Windowing**: dividing a continuous stream of data into discrete chunks (**windows**).
 - Count-based, Time-based, etc.



Windowing and Triggering

- ▶ **Windowing**: dividing a continuous stream of data into discrete chunks (**windows**).
 - Count-based, Time-based, etc.
- ▶ **Triggering**: defines **when** a computation **within a window** should be performed.



Windowing and Triggering

- ▶ **Windowing**: dividing a continuous stream of data into discrete chunks (**windows**).
 - Count-based, Time-based, etc.
- ▶ **Triggering**: defines **when** a computation **within a window** should be performed.
 - Count-based policy: the **maximum number** of tuples a window buffer can hold

Windowing and Triggering

- ▶ **Windowing:** dividing a continuous stream of data into discrete chunks (**windows**).
 - Count-based, Time-based, etc.

- ▶ **Triggering:** defines **when** a computation **within a window** should be performed.
 - Count-based policy: the **maximum number** of tuples a window buffer can hold
 - Time-based policy: based on **processing or event time** period



Windows After Triggering

- ▶ Two possibilities: tumbling and sliding

Windows After Triggering

- ▶ Two possibilities: tumbling and sliding
- ▶ Tumbling window: when the buffer fills up, all the tuples are evicted.

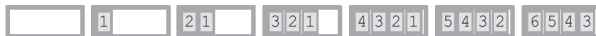


Windows After Triggering

- ▶ Two possibilities: **tumbling** and **sliding**
- ▶ **Tumbling window**: when the buffer fills up, **all** the tuples are **evicted**.

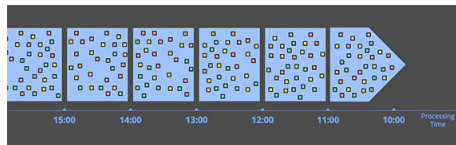


- ▶ **Sliding window**: when the buffer fills up, **older** tuples are **evicted**.



Triggering According to Processing Time

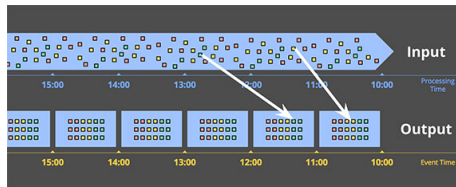
- ▶ The system **buffers up** incoming data into windows until **some amount of processing time has passed**.
- ▶ E.g., **five-minute** fixed windows



[<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>]

Triggering According to Event Time (1/2)

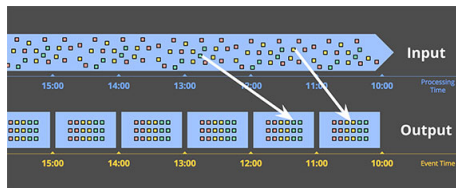
- Reflect the **times** at which **events** actually happened.



[<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>]

Triggering According to Event Time (1/2)

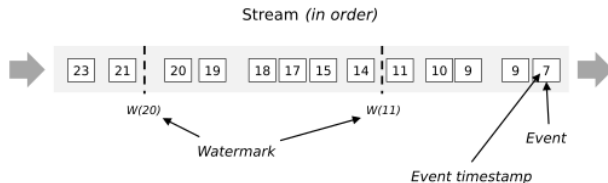
- ▶ Reflect the **times** at which **events** actually happened.
- ▶ Handling **out-of-order** events.



[<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>]

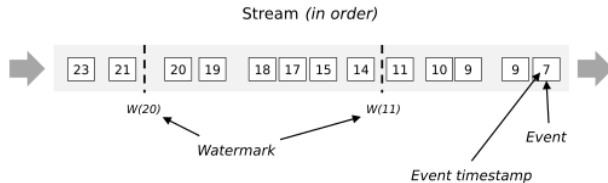
Triggering According to Event Time (2/2)

- **Watermarking** helps a stream processing system to deal with **lateness**.



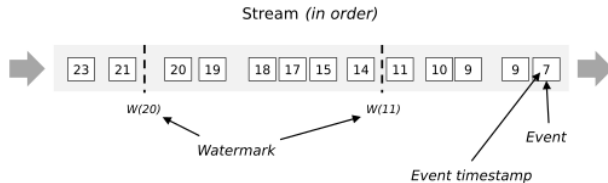
Triggering According to Event Time (2/2)

- ▶ **Watermarking** helps a stream processing system to deal with **lateness**.
- ▶ Watermarks **flow as part of the data stream** and carry a **timestamp t** .



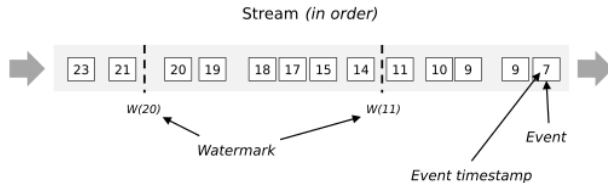
Triggering According to Event Time (2/2)

- ▶ **Watermarking** helps a stream processing system to deal with **lateness**.
- ▶ Watermarks **flow as part of the data stream** and carry a **timestamp t** .
- ▶ A $W(t)$ declares that **event time** has reached time t in that stream



Triggering According to Event Time (2/2)

- ▶ **Watermarking** helps a stream processing system to deal with **lateness**.
- ▶ Watermarks **flow as part of the data stream** and carry a **timestamp t** .
- ▶ A $W(t)$ declares that **event time** has reached time t in that stream
- ▶ A watermark is a **threshold** to specify **how long the system waits** for **late events**.



Spark Streaming

Spark Streaming

- Run a streaming computation as a **series** of very **small**, **deterministic batch jobs**.



Spark Streaming

- ▶ Run a streaming computation as a **series** of very **small**, **deterministic batch jobs**.
 - **Chops up** the live stream into batches of **X** seconds.
 - Treats each batch as **RDDs** and processes them using **RDD operations**.



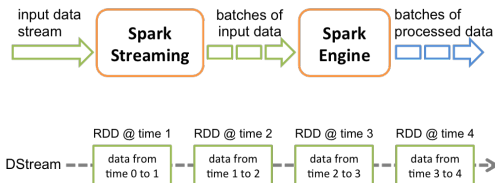
Spark Streaming

- ▶ Run a streaming computation as a **series** of very **small**, **deterministic batch jobs**.
 - **Chops up** the live stream into batches of **X** seconds.
 - Treats each batch as **RDDs** and processes them using **RDD operations**.
 - Discretized Stream Processing (**DStream**)



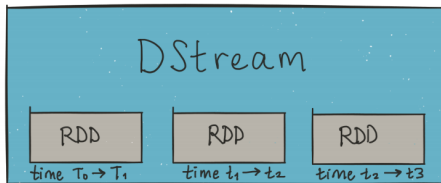
DStream (1/2)

- **DStream**: sequence of **RDDs** representing a stream of data.



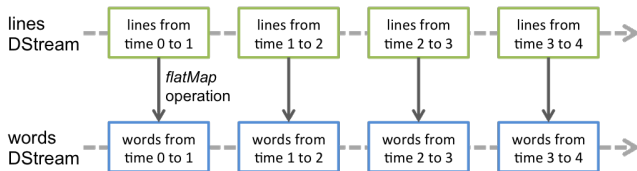
DStream (1/2)

- **DStream**: sequence of **RDDs** representing a stream of data.



DStream (2/2)

- Any **operation** applied on a **DStream** translates to operations on the underlying **RDDs**.





StreamingContext

- ▶ **StreamingContext** is the **main entry** point of all Spark Streaming functionality.

```
val conf = new SparkConf().setAppName(appName).setMaster(master)
val ssc = new StreamingContext(conf, Seconds(1))
```

- ▶ The second parameter, **Seconds(1)**, represents the **time interval** at which streaming data will be divided into **batches**.



Input Operations

► Socket connection

- Creates a DStream from text data received over a **TCP socket connection**.

```
ssc.socketTextStream("localhost", 9999)
```



Input Operations

► Socket connection

- Creates a DStream from text data received over a **TCP socket connection**.

```
ssc.socketTextStream("localhost", 9999)
```

► File stream

- Reads data from **files**.

```
streamingContext.fileStream[KeyClass, ValueClass, InputFormatClass](dataDirectory)  
streamingContext.textFileStream(dataDirectory)
```



Input Operations

► Socket connection

- Creates a DStream from text data received over a **TCP socket connection**.

```
ssc.socketTextStream("localhost", 9999)
```

► File stream

- Reads data from **files**.

```
streamingContext.fileStream[KeyClass, ValueClass, InputFormatClass](dataDirectory)  
streamingContext.textFileStream(dataDirectory)
```

► Connectors with external sources, e.g., **Twitter, Kafka, Flume, Kinesis, ...**



Transformations (1/2)

- ▶ Transformations on DStreams are still lazy!
- ▶ DStreams support many of the transformations available on normal Spark RDDs.



Transformations (1/2)

- ▶ Transformations on DStreams are still lazy!
- ▶ DStreams support many of the transformations available on normal Spark RDDs.
- ▶ Computation is kicked off explicitly by a call to the `start()` method.



Transformations (2/2)

- ▶ **map**: a new **DStream** by passing each **element** of the source DStream through a given function.



Transformations (2/2)

- ▶ **map**: a new **DStream** by passing each **element** of the source DStream through a given function.
- ▶ **reduce**: a new DStream of **single-element RDDs** by **aggregating** the elements in each RDD using a given function.



Transformations (2/2)

- ▶ **map**: a new **DStream** by passing each **element** of the source DStream through a given function.
- ▶ **reduce**: a new DStream of **single-element RDDs** by **aggregating** the elements in each RDD using a given function.
- ▶ **reduceByKey**: a new DStream of **(K, V) pairs** where the values for each key are **aggregated** using the given reduce function.



Example - Word Count (1/6)

- First we create a `StreamingContext`

```
import org.apache.spark._
import org.apache.spark.streaming._

// Create a local StreamingContext with two working threads and batch interval of 1 second.
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))
```



Example - Word Count (2/6)

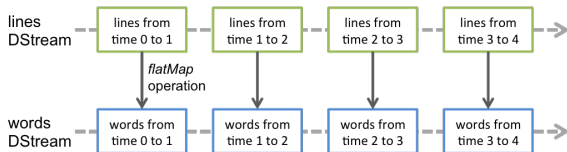
- ▶ Create a **DStream** that represents streaming data from a **TCP source**.
- ▶ Specified as **hostname** (e.g., localhost) and **port** (e.g., 9999).

```
val lines = ssc.socketTextStream("localhost", 9999)
```

Example - Word Count (3/6)

- ▶ Use `flatMap` on the stream to split the records text to words.
- ▶ It creates a new DStream.

```
val words = lines.flatMap(_.split(" "))
```





Example - Word Count (4/6)

- ▶ Map the **words** DStream to a DStream of **(word, 1)**.
- ▶ Get the **frequency of words** in each **batch of data**.
- ▶ Finally, **print** the result.

```
val pairs = words.map(word => (word, 1))  
val wordCounts = pairs.reduceByKey(_ + _)  
wordCounts.print()
```



Example - Word Count (5/6)

- ▶ Start the **computation** and **wait** for it to **terminate**.

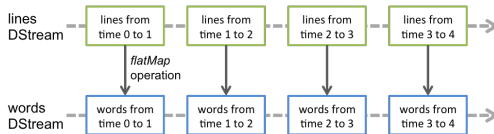
```
// Start the computation  
ssc.start()  
  
// Wait for the computation to terminate  
ssc.awaitTermination()
```

Example - Word Count (6/6)

```
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))

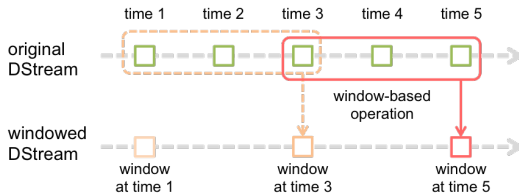
val lines = ssc.socketTextStream("localhost", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))
val wordCounts = pairs.reduceByKey(_ + _)
wordCounts.print()

ssc.start()
ssc.awaitTermination()
```



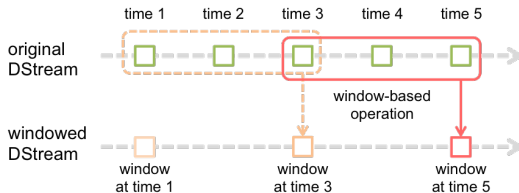
Window Operations (1/2)

- ▶ Spark provides a set of transformations that apply to a over a **sliding window** of data.



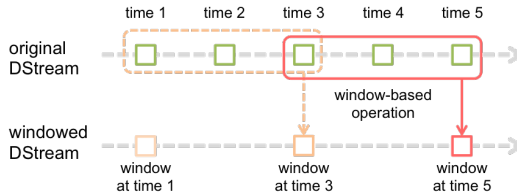
Window Operations (1/2)

- ▶ Spark provides a set of transformations that apply to a over a **sliding window** of data.
- ▶ A window is defined by two parameters: **window length** and **slide interval**.



Window Operations (1/2)

- ▶ Spark provides a set of transformations that apply to a over a **sliding window** of data.
- ▶ A window is defined by two parameters: **window length** and **slide interval**.
- ▶ A **tumbling window** effect can be achieved by making **slide interval = window length**





Window Operations (2/2)

- ▶ `reduceByWindow(func, windowLength, slideInterval)`
 - Returns a new **single-element DStream**, created by aggregating elements in the stream over a **sliding interval** using `func`.



Window Operations (2/2)

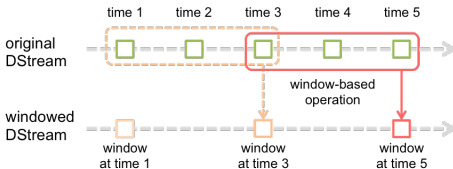
- ▶ `reduceByWindow(func, windowLength, slideInterval)`
 - Returns a new **single-element DStream**, created by aggregating elements in the stream over a **sliding interval** using `func`.
- ▶ `reduceByKeyAndWindow(func, windowLength, slideInterval)`
 - Called on a DStream of **(K, V) pairs**.
 - Returns a **new DStream of (K, V) pairs** where the values for each key are aggregated using function `func` over **batches in a sliding window**.

Example - Word Count with Window

```
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))

val lines = ssc.socketTextStream("localhost", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))
val windowedWordCounts = pairs.reduceByKeyAndWindow(_ + _, Seconds(30), Seconds(10))
windowedWordCounts.print()

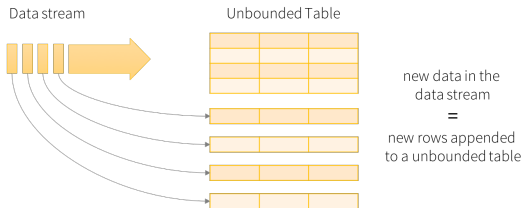
ssc.start()
ssc.awaitTermination()
```



Structured Streaming

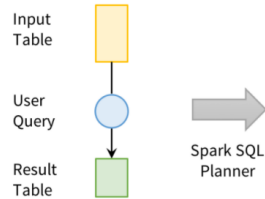
Structured Streaming

- Treating a **live data stream** as a **table** that is being **continuously appended**.

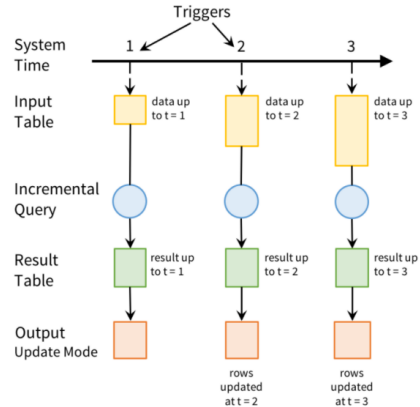


Data stream as an unbounded table

Programming Model



User's batch-like
query on input table



Incremental execution on streaming data



Output Modes

► Three output modes:

1. **Append**: only the new rows **appended to the result table** since the last trigger will be written to the external storage.



Output Modes

► Three output modes:

1. **Append**: only the new rows **appended to the result table** since the last trigger will be written to the external storage.
2. **Complete**: the **entire updated result table** will be written to external storage.

► Three output modes:

1. **Append**: only the new rows **appended to the result table** since the last trigger will be written to the external storage.
2. **Complete**: the **entire updated result table** will be written to external storage.
3. **Update**: only the rows that were **updated in the result table** since the last trigger will be changed in the external storage.



Steps to Define a Streaming Query (1/4)

- ▶ Define **input sources**.
- ▶ Use `spark.readStream` to create a `DataStreamReader`.

```
val spark = SparkSession.builder.master("local[2]").appName("appname").getOrCreate()

val lines = spark.readStream.format("socket")
    .option("host", "localhost")
    .option("port", 9999)
    .load()
```



Steps to Define a Streaming Query (2/4)

- ▶ Transform data.
- ▶ E.g., below `counts` is a `streaming DataFrame` that represents the running word counts.

```
import org.apache.spark.sql.functions._  
  
val words = lines.select(split(col("value"), " ").as("word"))  
  
val wordCounts = words.groupBy("word").count()
```

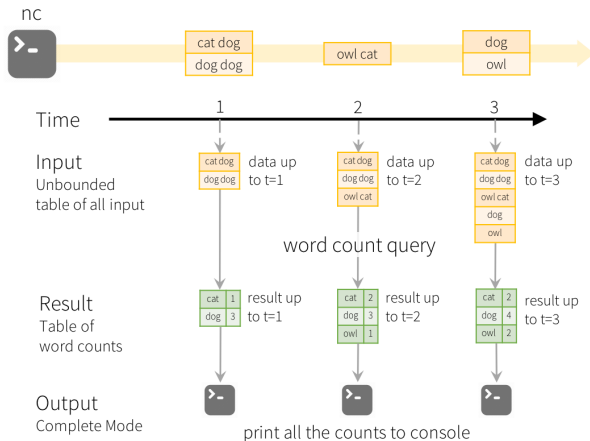


Steps to Define a Streaming Query (3/4)

- ▶ Define **output sink** and **output mode**.
- ▶ Use `DataFrame.writeStream` to define how to write the processed output data.
- ▶ **Start** the query.

```
val query = wordCounts.writeStream.format("console").outputMode("complete").start()  
query.awaitTermination()
```

Steps to Define a Streaming Query (4/4)



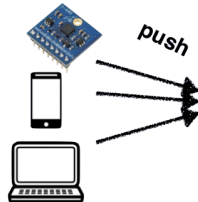
[<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>]

Data Stream Storage

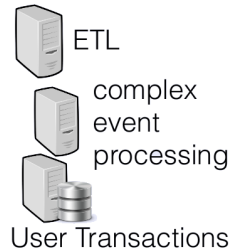
The Problem

- ▶ We need disseminate streams of events from various producers to various consumers.

Data Producers

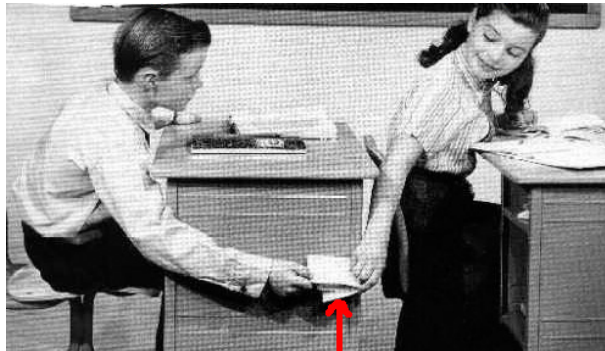


Data Consumers



Possible Solution?

- ▶ Messaging systems



Message

www.defit.org



What is Messaging System?

- ▶ **Messaging system** is an approach to **notify consumers** about new events.

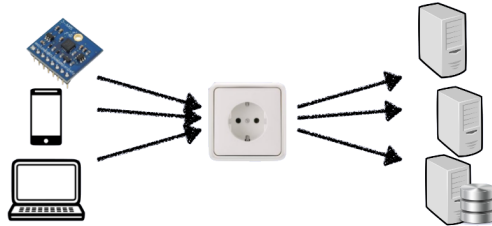


What is Messaging System?

- ▶ **Messaging system** is an approach to **notify consumers** about new events.
- ▶ **Messaging systems**
 - **Direct** messaging
 - Message **brokers**

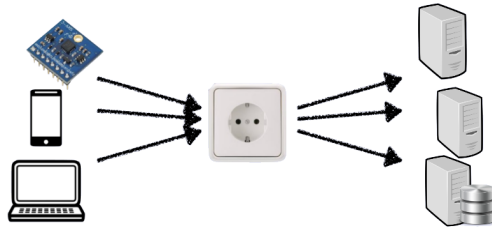
Direct Messaging (1/2)

- ▶ Necessary in **latency critical** applications (e.g., remote surgery).
- ▶ A **producer** sends a message containing the event, which is **pushed** to **consumers**.



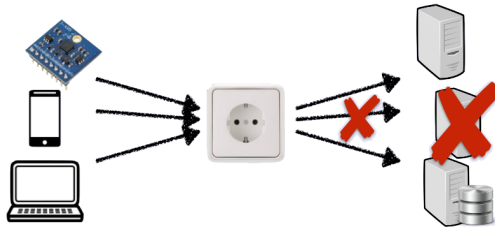
Direct Messaging (1/2)

- ▶ Necessary in **latency critical** applications (e.g., remote surgery).
- ▶ A **producer** sends a message containing the event, which is **pushed** to **consumers**.
- ▶ Both consumers and producers have to be **online at the same time**.



Direct Messaging (2/2)

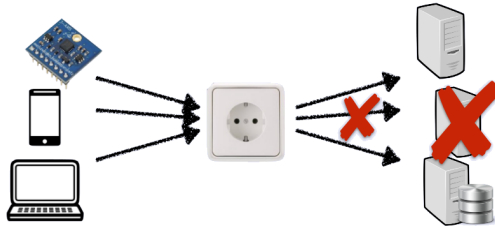
- What happens if a **consumer** **crashes** or temporarily **goes offline**? (**not durable**)



-
- The diagram shows a central square box with a circular port, representing a processing unit. On the left, four input devices (a blue circuit board, a smartphone, and a laptop) have arrows pointing towards the central box. On the right, three output devices (a server tower, a monitor, and a server rack with a disk) have arrows pointing away from the central box. A large red 'X' is superimposed over the output devices, indicating a failure or error in the output process.

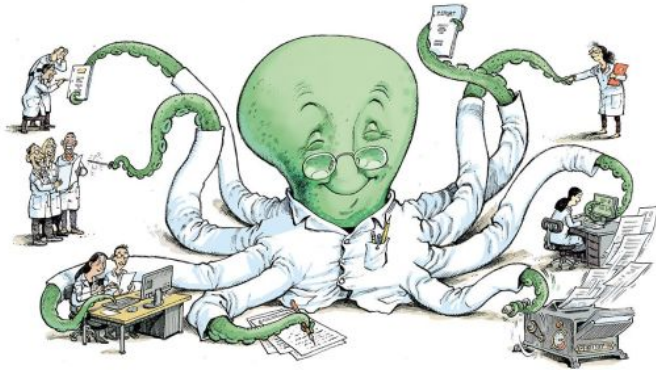
Direct Messaging (2/2)

- ▶ What happens if a **consumer crashes** or temporarily **goes offline**? (**not durable**)
- ▶ What happens if **producers** send messages **faster** than the **consumers** can process?
 - **Dropping** messages
 - **Backpressure**



-
- The diagram shows three input devices on the left: a blue circuit board, a black smartphone, and a laptop. Arrows from these devices converge on a central white square switch with a circular port. From the switch, three arrows point to a stack of three server units on the right. A large red 'X' is superimposed over the arrows and the servers, indicating a failure or error in the connection.

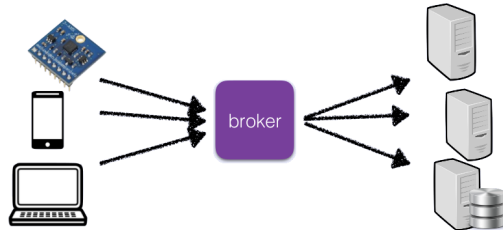
Message Broker



[<https://bluesyemre.com/2018/10/16/thousands-of-scientists-publish-a-paper-every-five-days>]

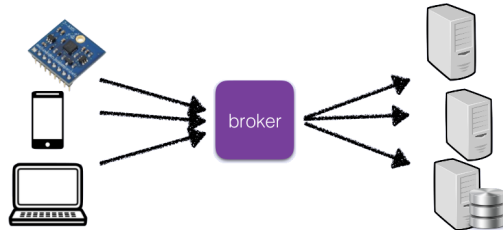
Message Broker

- ▶ A **message broker** decouples the **producer-consumer** interaction.
- ▶ It runs as a **server**, with **producers and consumers** connecting to it as **clients**.



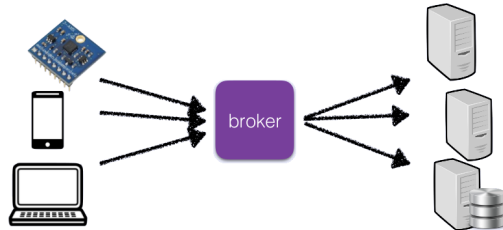
Message Broker

- ▶ A **message broker** decouples the **producer-consumer** interaction.
- ▶ It runs as a **server**, with **producers and consumers** connecting to it as **clients**.
- ▶ **Producers** write messages to the broker, and **consumers** receive them by reading them from the broker.



Message Broker

- ▶ A **message broker** decouples the **producer-consumer** interaction.
- ▶ It runs as a **server**, with **producers and consumers** connecting to it as **clients**.
- ▶ **Producers** write messages to the broker, and **consumers** receive them by reading them from the broker.
- ▶ **Consumers** are generally **asynchronous**.





Logs-Based Message Broker

- ▶ In typical message brokers, once a message is **consumed**, it is **deleted**.



Logs-Based Message Broker

- ▶ In typical message brokers, once a message is **consumed**, it is **deleted**.
- ▶ **Log-based message brokers** **durably** store all events in a sequential **log**.



Logs-Based Message Broker

- ▶ In typical message brokers, once a message is **consumed**, it is **deleted**.
- ▶ **Log-based message brokers** **durably** store all events in a sequential **log**.
- ▶ A **log** is an **append-only** sequence of records on **disk**.



Logs-Based Message Broker

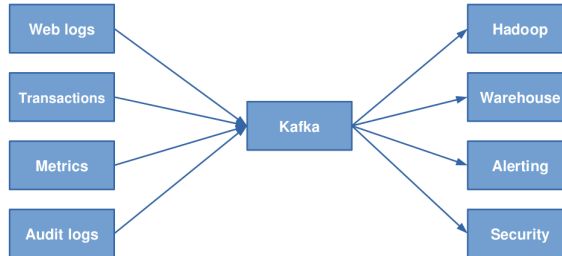
- ▶ In typical message brokers, once a message is **consumed**, it is **deleted**.
- ▶ **Log-based message brokers** **durably** store all events in a sequential **log**.
- ▶ A **log** is an **append-only** sequence of records on **disk**.
- ▶ A **producer** sends a message by **appending** it to the end of the log.
- ▶ A **consumer** receives messages by reading the log **sequentially**.

Kafka - A Log-Based Message Broker



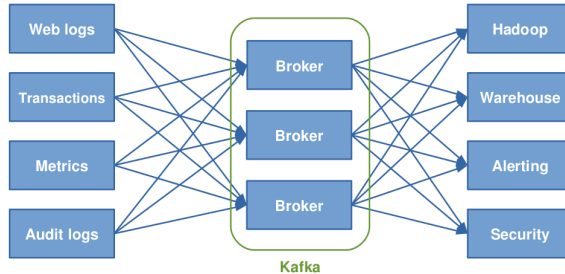
Kafka (1/5)

- **Kafka** is a distributed, topic oriented, partitioned, replicated commit **log service**.



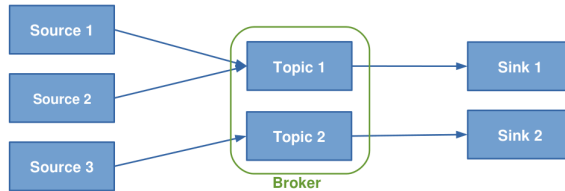
Kafka (2/5)

- Kafka is a **distributed**, topic oriented, partitioned, replicated commit **log service**.



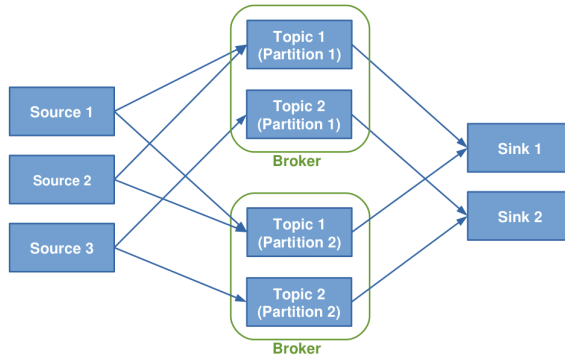
Kafka (3/5)

- Kafka is a distributed, topic oriented, partitioned, replicated commit log service.



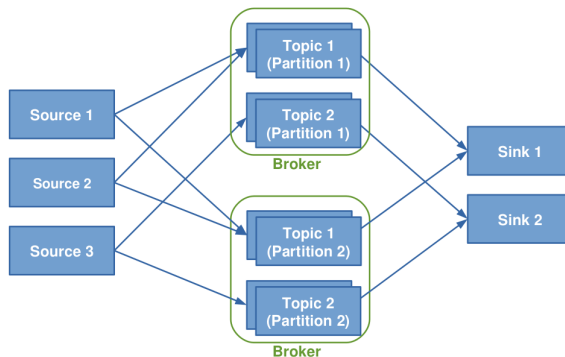
Kafka (4/5)

- Kafka is a distributed, topic oriented, partitioned, replicated commit log service.



Kafka (5/5)

- Kafka is a distributed, topic oriented, partitioned, replicated commit log service.



Topics and Partition (1/6)

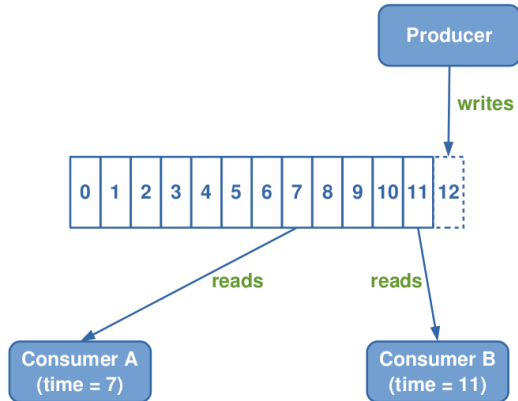
- **Topics** are **queues**: a **stream of messages** of a **particular type**

```
jkreps-mn:~ jkreps$ tail -f -n 20 /var/log/apache2/access_log
::1 - - [23/Mar/2014:15:07:00 -0700] "GET /images/apache_feather.gif HTTP/1.1" 200 4128
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/producer_consumer.png HTTP/1.1" 200 86
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/log_anatomy.png HTTP/1.1" 200 19579
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/consumer-groups.png HTTP/1.1" 200 268
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/log_compaction.png HTTP/1.1" 200 4141
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /documentation.html HTTP/1.1" 200 189893
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/log_cleaner_anatomy.png HTTP/1.1" 200
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/kafka_log.png HTTP/1.1" 200 134321
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/mirror-maker.png HTTP/1.1" 200 17054
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /documentation.html HTTP/1.1" 200 189937
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /styles.css HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/kafka_logo.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/producer_consumer.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/log_anatomy.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/consumer-groups.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/log_cleaner_anatomy.png HTTP/1.1" 304
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/log_compaction.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/kafka_log.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/mirror-maker.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:09:55 -0700] "GET /documentation.html HTTP/1.1" 200 195264
```



Topics and Partition (2/6)

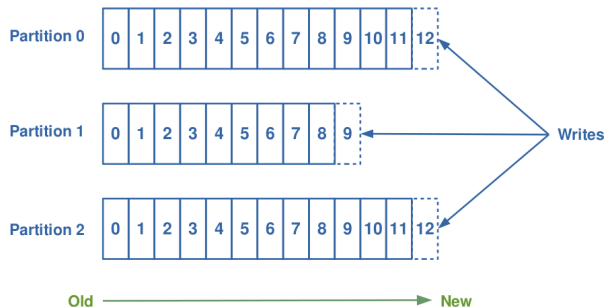
- Each **message** is assigned a **sequential id** called an **offset**.



Topics and Partition (3/6)

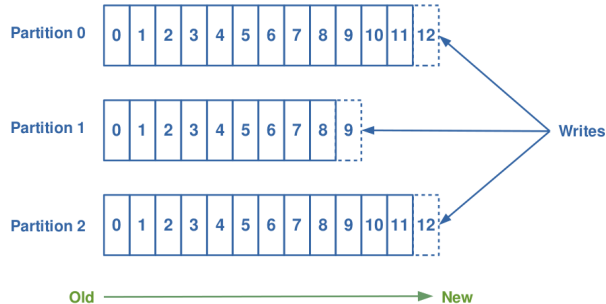
► Topics are **logical** collections of **partitions** (the **physical files**).

- Ordered
- Append only
- Immutable



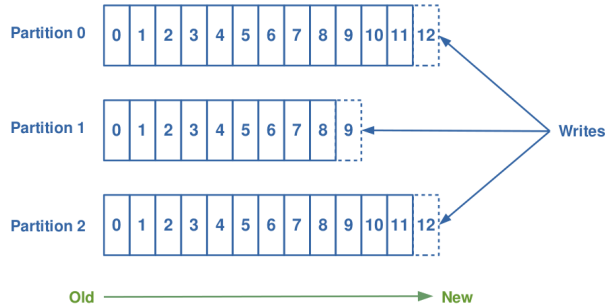
Topics and Partition (4/6)

- Ordering is only guaranteed **within** a partition for a topic.



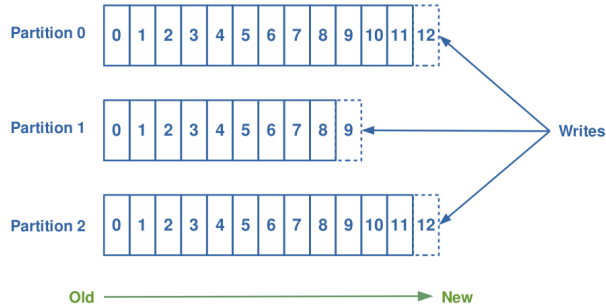
Topics and Partition (4/6)

- ▶ Ordering is only **guaranteed within** a **partition** for a **topic**.
- ▶ Messages sent by a **producer** to a particular topic partition will be **appended** in the order they are sent.



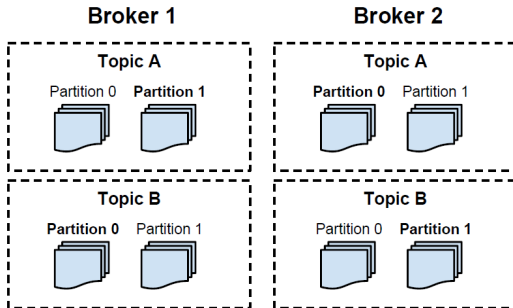
Topics and Partition (4/6)

- ▶ Ordering is only **guaranteed within** a **partition** for a **topic**.
- ▶ Messages sent by a **producer** to a particular topic partition will be **appended** in the order they are sent.
- ▶ A **consumer** instance sees messages in the order they are stored in the log.



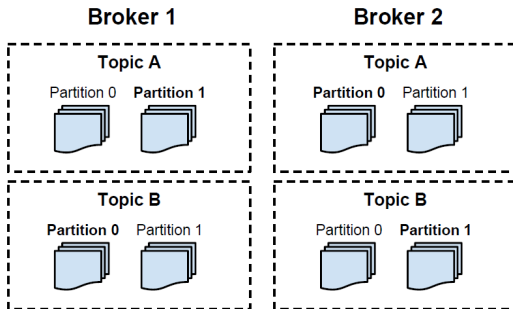
Topics and Partition (5/6)

- Partitions of a topic are replicated: fault-tolerance



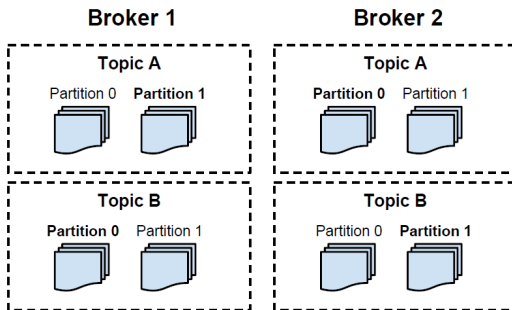
Topics and Partition (5/6)

- ▶ **Partitions** of a topic are **replicated**: **fault-tolerance**
- ▶ A **broker** contains some of the **partitions** for a topic.

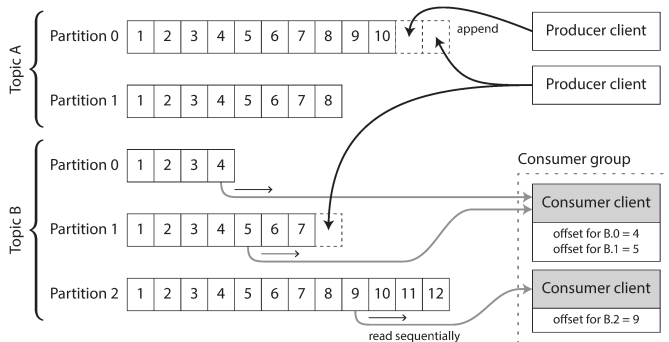


Topics and Partition (5/6)

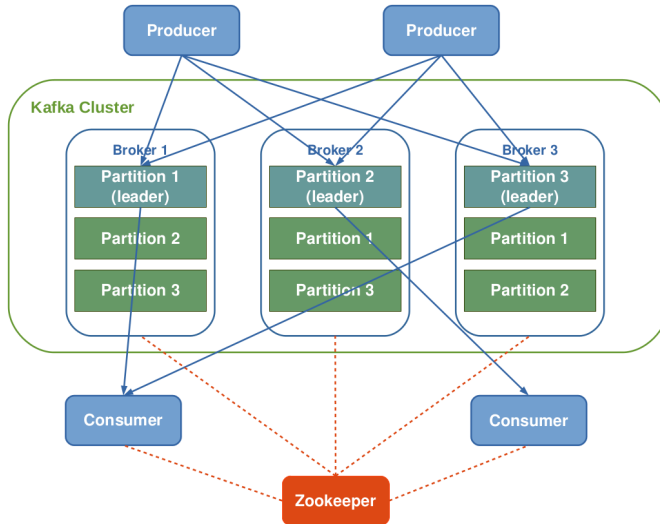
- ▶ **Partitions** of a topic are **replicated**: **fault-tolerance**
- ▶ A **broker** contains some of the **partitions** for a topic.
- ▶ One broker is the **leader** of a partition: all **writes** and **reads** must go to the leader.



Topics and Partition (6/6)



Kafka Architecture



Summary



Summary

- ▶ Data stream, unbounded data, tuples
- ▶ Event-time vs. processing time
- ▶ Windowing and triggering
- ▶ Messaging system and partitioned logs
- ▶ Kafka: distributed, topic oriented, partitioned, replicated log service
- ▶ Spark streaming and structured streaming

- ▶ J. Kreps et al., “Kafka: A distributed messaging system for log processing”, NetDB 2011
- ▶ M. Zaharia et al., “Spark: The Definitive Guide”, O'Reilly Media, 2018 - Chapter 20
- ▶ T. Akidau et al., “The dataflow model: a practical approach to balancing correctness, latency, and cost in massive-scale, unbounded, out-of-order data processing”, VLDB 2015.
- ▶ M. Fragkoulis et al., “A Survey on the Evolution of Stream Processing Systems”, 2020
- ▶ T. Akidau, “The world beyond batch: Streaming 101”,
<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>

Questions?