

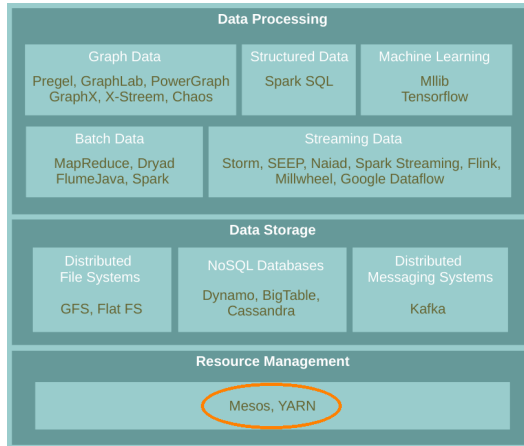


Resource Management - Mesos, YARN, and Borg

Amir H. Payberah
payberah@kth.se
2025-09-23



Where Are We?





Motivation

- ▶ Rapid innovation in cloud computing.
- ▶ No single framework optimal for all applications.
- ▶ Running each framework on its dedicated cluster:
 - Expensive
 - Hard to share data



Proposed Solution

- ▶ Running **multiple frameworks** on a **single cluster**.
- ▶ Maximize **utilization** and **share** data between frameworks.
- ▶ Three resource management systems:
 - Mesos
 - YARN
 - Borg

Question?

How to **schedule** resource offering among **frameworks**?



Schedule Frameworks

- ▶ Monolithic scheduler
- ▶ Two-Level scheduler

Monolithic Scheduler (1/2)

▶ Job requirements


- Response time
- Throughput
- Availability

▶ Job execution plan

- Task DAG
- Inputs/outputs

▶ Estimates

- Task duration
- Input sizes
- Transfer sizes





Monolithic Scheduler (2/2)


► Advantages

- Can achieve **optimal** schedule.

► Disadvantages

- **Complexity**: hard to scale and ensure resilience.
- Hard to anticipate **future frameworks** requirements.
- Need to **refactor** existing frameworks.

Two-Level Scheduler (1/2)





Two-Level Scheduler (2/2)

► Advantages


- **Simple**: easier to scale and make resilient.
- **Easy to port** existing frameworks, support new ones.

► Disadvantages

- Distributed scheduling decision: **not optimal**.

Two-Level vs. Monolithic

- ▶ **Two-level schedulers:** separate concerns of **resource allocation** and **task placement**.
 - An active **resource manager** offers **compute resources** to multiple parallel, independent **scheduler frameworks**.
 - Mesos and Yarn
- ▶ **Monolithic schedulers:** use a single, **centralized scheduling** algorithm for **all jobs**.
 - Borg



[Schwarzkopf et al., Omega: flexible, scalable schedulers for large compute clusters, EuroSys'13.]


Mesos

- **Mesos** is a common **resource sharing** layer, over which diverse frameworks can run.




Computation Model

- ▶ A **framework** (e.g., Hadoop, Spark) manages and runs one or more **jobs**.
- ▶ A **job** consists of one or more **tasks**.
- ▶ A **task** (e.g., map, reduce) consists of one or more **processes** running on same machine.



Fine-Grained Sharing

- Allocation at the level of **tasks** within a **job**.

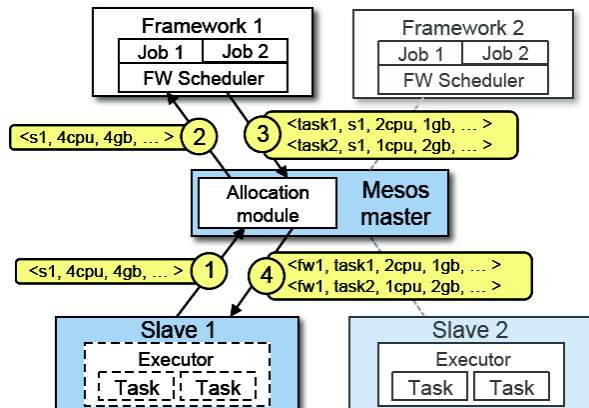




Mesos Scheduler


- ▶ Master sends resource offers to frameworks.
- ▶ Frameworks select which offers to accept and which tasks to run.
- ▶ Unit of allocation: resource offer
 - Vector of available resources on a node
 - For example, node1: $\langle 1\text{CPU}, 1\text{GB} \rangle$, node2: $\langle 4\text{CPU}, 16\text{GB} \rangle$

Mesos Architecture (1/4)




- Slaves continuously send status updates about resources to the Master.

Mesos Architecture (2/4)



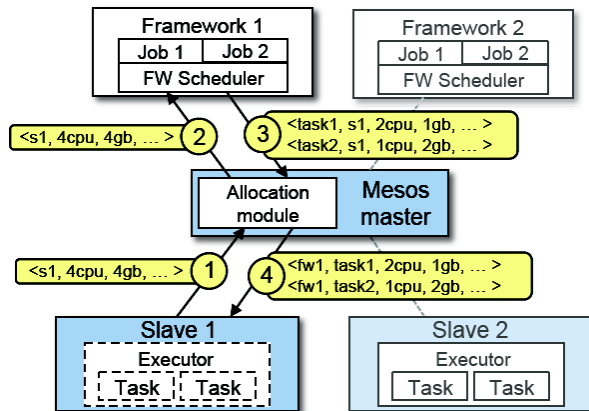
- Pluggable **scheduler** picks **framework** to send an **offer** to.

Mesos Architecture (3/4)



- Framework scheduler selects resources and provides tasks.

Mesos Architecture (4/4)



- Framework executors launch tasks.

Question?

How to allocate resources of different types?

-



Max-Min Fairness - Example

- ▶ 1 resource: CPU
- ▶ Total resources: 20 CPU
- ▶ User 1 has x tasks and wants $\langle 1\text{CPU} \rangle$ per task
- ▶ User 2 has y tasks and wants $\langle 2\text{CPU} \rangle$ per task

$\max(x, y)$ (maximize allocation)

subject to

$x + 2y \leq 20$ (CPU constraint)

$x = 2y$

so

$x = 10$

$y = 5$



Properties of Max-Min Fairness

► Share guarantee

- Each user can get at least $\frac{1}{n}$ of the resource.
- But will get less if her demand is less.

► Strategy proof

- Users are not better off by asking for more than they need.
- Users have no reason to lie.


Question?

When is Max-Min Fairness **NOT** Enough?


Need to schedule **multiple, heterogeneous** resources, e.g.,
CPU, memory, etc.

Problem

- 1 resource: CPU
- User 1 wants $\langle 1\text{CPU} \rangle$ per task
- User 2 wants $\langle 2\text{CPU} \rangle$ per task



- 2 resources: CPUs and mem
- User 1 wants $\langle 1\text{CPU}, 4\text{GB} \rangle$ per task
- User 2 wants $\langle 2\text{CPU}, 1\text{GB} \rangle$ per task
- What is a fair allocation?



A Natural Policy (1/2)

- ▶ **Asset fairness**: give weights to resources (e.g., 1 CPU = 1 GB) and **equalize total value given to each user**.
- ▶ Total resources: 28 CPU and 56GB RAM (e.g., 1 CPU = 2 GB)
 - User 1 has x tasks and wants $\langle 1\text{CPU}, 2\text{GB} \rangle$ per task
 - User 2 has y tasks and wants $\langle 1\text{CPU}, 4\text{GB} \rangle$ per task
- ▶ Asset fairness yields:

$$\max(x, y)$$


$$x + y \leq 28$$


$$2x + 4y \leq 56$$

$$2x = 3y$$

User 1: $x = 12$: $\langle 43\%\text{CPU}, 43\%\text{GB} \rangle$ ($\sum = 86\%$)

User 2: $y = 8$: $\langle 28\%\text{CPU}, 57\%\text{GB} \rangle$ ($\sum = 86\%$)





- ▶ **Problem:** violates share grantee.
- ▶ User 1 gets less than 50% of both CPU and RAM.
- ▶ Better off in a separate cluster with half the resources.



Challenge

- ▶ Can we find a fair sharing policy that provides:
 - Share guarantee
 - Strategy-proofness

- ▶ Can we generalize max-min fairness to multiple resources?



Proposed Solution

Dominant Resource Fairness (**DRF**)

Dominant Resource Fairness (DRF) (1/2)

- ▶ **Dominant resource** of a user: the resource that user has the **biggest share of**.
 - Total resources: $\langle 8\text{CPU}, 5\text{GB} \rangle$
 - User 1 allocation: $\langle 2\text{CPU}, 1\text{GB} \rangle$: $\frac{2}{8} = 25\%$ CPU and $\frac{1}{5} = 20\%$ RAM
 - Dominant resource of User 1 is **CPU** ($25\% > 20\%$)
- ▶ **Dominant share** of a user: the **fraction** of the **dominant resource** she is allocated.
 - User 1 dominant share is **25%**.

Dominant Resource Fairness (DRF) (2/2)

- ▶ Apply **max-min fairness** to **dominant shares**: give every user an equal share of her dominant resource.
- ▶ **Equalize** the **dominant share** of the users.
 - Total resources: $\langle 9\text{CPU}, 18\text{GB} \rangle$
 - User 1 wants $\langle 1\text{CPU}, 4\text{GB} \rangle$; Dominant resource: RAM ($\frac{1}{9} < \frac{4}{18}$)
 - User 2 wants $\langle 3\text{CPU}, 1\text{GB} \rangle$; Dominant resource: CPU ($\frac{3}{9} > \frac{1}{18}$)
- ▶ $\max(x, y)$


$$x + 3y \leq 9$$

$$4x + y \leq 18$$

$$\frac{4x}{18} = \frac{3y}{9}$$

User 1: $x = 3$: $\langle 33\%\text{CPU}, 66\%\text{GB} \rangle$

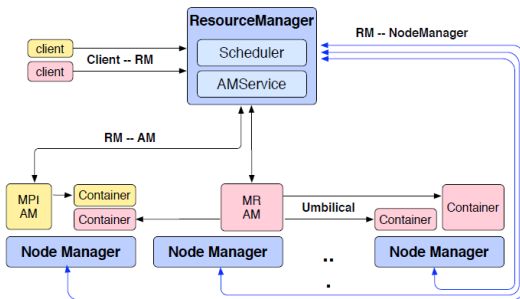
User 2: $y = 2$: $\langle 66\%\text{CPU}, 16\%\text{GB} \rangle$



YARN


YARN Architecture

- ▶ Resource Manager (RM)
- ▶ Application Master (AM)
- ▶ Node Manager (NM)




YARN Architecture - Resource Manager

- ▶ One per cluster (Central: global view)
- ▶ Job requests are submitted to RM.
 - To start a job, RM finds a container to spawn AM.
- ▶ Only handles an overall resource profile for each job.
 - Local optimization is up to the job.




YARN Architecture - Application Manager

- ▶ The head of a job.
- ▶ Runs as a container.
- ▶ Request resources from RM (num. of containers/resource per container/locality ...)



YARN Architecture - Node Manager

- ▶ The **worker daemon**.
- ▶ Registers with RM.
- ▶ **One** per node.
- ▶ **Report resources** to RM: memory, CPU, ...



Borg



Borg


- ▶ Cluster management system at Google.



-
- The diagram illustrates a Cell architecture. At the top, a light blue box labeled "Job" is enclosed within a larger light blue box labeled "Alloc set". Below the "Alloc set", there are two light blue boxes, each labeled "Task". Arrows point from the "Alloc set" box to each of the "Task" boxes. To the left of the "Task" boxes, the text "Alloc instance" is written. Below the "Task" boxes, there is a horizontal row of four server nodes, each represented by a light blue box with a circular icon and a vertical bar. Arrows point from each of the "Task" boxes to one of the server nodes. The entire diagram is set against a light blue background.


Borg Architecture

- ▶ BorgMaster
 - The **central brain** of the system
 - Holds the **cluster state**
 - **Replicated** for **reliability** (using paxos)
 - **Scheduling**: where to **place tasks**?
- ▶ Borglet
 - **Manage and monitor** tasks and resource
 - BorgMaster **polls Borglet** every few seconds




Borg Scheduler

- ▶ Feasibility checking: find machines for a given job
- ▶ Scoring: pick one machines
- ▶ According to the users prefs and built-in criteria



Docker and Kubernetes


Application Deployment





Traditional Deployment Era


- ▶ Running applications on **physical servers**.
- ▶ **No resource boundaries** for applications in a physical server
- ▶ **Resource allocation** issues, e.g., one application would take up most of the resources, so the other applications would underperform.



Traditional Deployment

Virtualized Deployment Era


- ▶ **Virtual Machines (VMs)**: a **full machine** running all the components, including its own operating system (OS), on top of the **virtualized hardware**.
- ▶ Virtualization allows to run **multiple VMs** on a **single physical server's CPU**.
 - **Utilizes the resources** of a physical server better.
 - Better **scalability** as applications can be **added/updated** easily.



Virtualized Deployment



Container Deployment Era

- ▶ **Containers** are similar to **VMs**, but they have **relaxed isolation** properties to **share the OS** among the applications.
- ▶ Similar to a **VM**, a **container** packages applications as **images** that contain **everything needed to run them**: code, runtime environment, libraries, and configuration.
- ▶ As they are **decoupled** from the **underlying infrastructure**, they are **portable** across clouds and OS distributions.



Container Deployment

- ▶ Docker is a virtualization software.
- ▶ A docker image is a template, and a container is a copy of that template.





Container Orchestration

- ▶ Container **scalability** is an **operational challenge**.
- ▶ If we have **10 containers** and four applications, it is **not difficult** to manage the **deployment and maintenance** of the containers.
- ▶ But, what if we have **1000 containers** and **400 services**?
- ▶ Container **orchestration** can help to **manage the lifecycles of containers**, especially in large and dynamic environments.
- ▶ Container orchestration tools: **Kubernetes** (based on Borg), **Marathon** (runs on Mesos)





Questions

- ▶ **Who** gets to decide weights in weighted max-min fairness and **what** biases might this introduce?

Possible Answers

- ▶ **System admins:** may favor powerful users or paying customers.
- ▶ **Organizational priorities:** research areas with more funding may get higher weights.
- ▶ **Bias risk:** dominant groups reinforce their advantage, marginalized groups get fewer resources.
- ▶ **Opaque choices:** if weight rules aren't transparent, users can't contest unfair allocations.
- ▶ **Equity gap:** "fair" weights may ignore social context (e.g., small labs, NGOs need proportionally more).



Questions

- ▶ Fairness in clusters often means equal technical access. How could we design systems that account for **social context**, e.g., prioritizing **under-resourced groups**?



Possible Answers

- ▶ Allocate **extra resources** to groups with **fewer starting advantages**.
- ▶ Scheduling that considers **deadlines**, **social impact**, or **community benefit**, not just efficiency.
- ▶ Let **affected groups** help set fairness rules, instead of only system admins.
- ▶ Show **who** got resources, **why**, and with **what** effects.

Summary

- ▶ Mesos
 - Offered-based
 - Max-Min fairness: DRF
- ▶ YARN
 - Request-based
 - RM, AM, NM
- ▶ Borg
 - Request-based
 - BorgMaster, Borglet
 - Kubernetes

- ▶ B. Hindman et al., “Mesos: A Platform for Fine-Grained Resource Sharing in the Data Center”, NSDI 2011
- ▶ V. Vavilapalli et al., “Apache hadoop yarn: Yet another resource negotiator”, ACM Cloud Computing 2013
- ▶ A. Verma et al., “Large-scale cluster management at Google with Borg”, EuroSys 2015

Questions?

Acknowledgements

Some slides were derived from Ion Stoica and Ali Ghodsi slides (Berkeley University), Wei-Chiu Chuang slides (Purdue University), and Arnon Rotem-Gal-Oz (Amdocs).